

Predictive Analytics Lecture 5

Adam Kapelner

Stat 422/722

at The Wharton School of the University of Pennsylvania

February 14 & 15, 2017


Underfitting & Overfitting

$$Y = f(x_1, \dots, x_p) + \mathcal{E}$$

Goal of machine learning: fit f as best as possible. When we build models, we do one (or both of the following)

- We fall **short** by underfitting (usually due to too little degrees of freedom and inflexible bases). For example: if the f is a curve and we fit a line, we underfit (recall medicorp sales vs bonus regression).
- We can shoot too **long** by encroaching on and fitting / optimizing to the \mathcal{E} . Since \mathcal{E} is independent of x_1, \dots, x_p , this part of \hat{f} is essentially a random fit and it is the opposite of the “data-driven approach”.

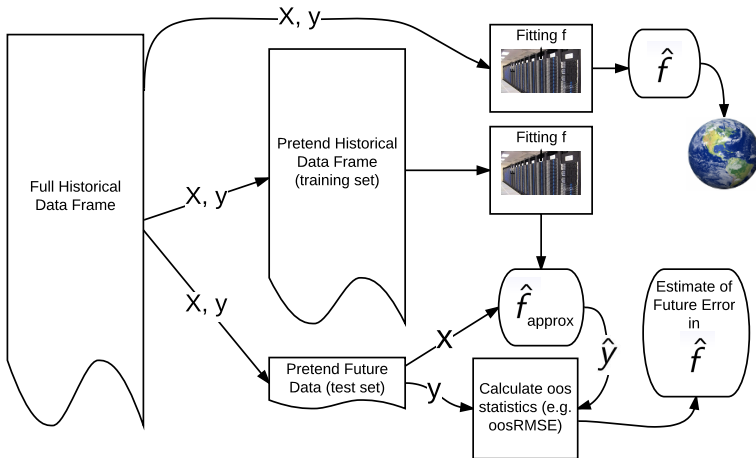
Complexity-Fit Tradeoff



Blue is in-sample fit metric and red is oos fit metric. This is Fig 7.1 from Hastie and Tibsharani (2009).

Assessment: OOS Validation

But knowing where you are on that y-axis would involve knowing the truth. We need to estimate this, so we use oos validation:



Assumptions and Tradeoffs when Splitting

We have a choice to split our dataframe into two pieces. Assuming each data point is independent (the running assumption), you should do this completely randomly. When would this assumption not be true? For example, a time series.

Additionally, we need to assume a non-stationary model relationship. So,

$$Y = f(x_1, \dots, x_p) + \mathcal{E} \quad \text{and not} \quad Y = f_t(x_1, \dots, x_p) + \mathcal{E}$$

where f changes with time. In essence non-stationarity is a lack of generalization and when predicting, it is a form of extrapolation.

How large should the test set be? Usual sizes are 10-30%. What's the tradeoff? If the test set is larger, then ...

- 1 the more accurate the assessment of generalization error would be (less variance) and
- 2 the less accurate the model will be since it's fitting with less data (more bias)

If the test set is smaller then, vice versa. Note: the in-sample and oos statistics are statistics! Thus, they are random!


Less Randomness in the OOS Statistics

If we change the observations in the training/test splits, we will get different models and different estimates of future error. Thus, our oosRMSE was really oosRMSE conditional on the idiosyncratic split we happened to get!

We can at the very least... get rid of this idiosyncratic error by ... averaging over all training-test splits. If we have $n = 100$ and the test set is 10%, that means we only have $\binom{100}{10} = 1.73 \times 10^{13}$ split configurations to average over!

We can approximate the averaging over all splits by just taking $\frac{100\%}{10\%} = 10$ random but unique splits called **folds**. Thus, each observation is represented in the test set once (leading to a more stable estimate). This is known as **K-fold cross validation (CV)** where here $K = \frac{100\%}{10\%} = 10$ (and this procedure seems to be the industry standard).

10-fold CV




$K = 10$ is arbitrary (but remember where it was based on: the 10-30% test set recommendation). In practice, I've only used 5 or 10 fold CV.

This does not really solve any of our big problems but gives us a little boost in terms of a reduction in standard error of our generalization error estimate. That's OK; we can take all the help we can get if it's costless!!

Note 1: If $K = n$ then we use the test set as one sample; this is known as "leave one out CV" (LOOCV) and it is not recommended — high variance and computational cost! Note 2: bleeding edge of stats — find CI's for generalization error.

$K = 3$ -fold CV on a Linear Model Ex. 1/5


Imagine the following data $n = 9$ where we are fitting a response by one feature (ignore the colors):



Imagine we choose a linear model.

$K = 3$ -fold CV on a Linear Model Ex. 2/5

In the first fold, the red is left out and thus we fit a line to the blue and green points:




Then we calculate the residuals to the red points (the test set in this fold) and calculate

$$SSE = (2 - 2.2)^2 + (3.8 - 2.1)^2 + (3.5 - 2.05)^2 = 5.03$$

$K = 3$ -fold CV on a Linear Model Ex. 3/5

In the second fold, the green is left out and thus we fit a line to the blue and red points:




Then we calculate the residuals to the red points (the test set in this fold) and calculate

$$SSE = (1.2 - 2.3)^2 + (3.4 - 2.25)^2 + (1.3 - 2.2)^2 = 3.34$$

$K = 3$ -fold CV on a Linear Model Ex. 4/5

In the third (last) fold, the blue is left out and thus we fit a line to the green and red points:



Then we calculate the residuals to the red points (the test set in this fold) and calculate

$$SSE = (2.35 - 2.3)^2 + (2.4 - 1.5)^2 + (2.2 - 1.4)^2 = 1.45$$

$K = 3$ -fold CV on a Linear Model Ex. 1/5

Then we aggregate all oos results together (SSE's are additive) and we can compute a final oos statistics e.g. the oosRMSE:

$$oosRMSE = \sqrt{\frac{5.03 + 3.34 + 1.45}{9}} = 1.045$$

Limits of JMP / Intro to R's MLR Package

And... JMP can't do K -fold CV! (Except in one limited case which doesn't help us right now). But of course R can do it... [R Demo with MLR]

Validating Multiple Models

Let's look at a few models for the White Wine data with no validation (but no cross-validation). Here the response is wine quality as measured by professional raters and features are 11 features (e.g. acidity, sugar, pH and alcohol content).

- A plain linear model
- B 6-degree polynomials for all features
- C 6° polynomials and all 1st order interactions
- D 6° polynomials and all interactions up to 3rd degree (AKA 2nd order)
- E 6° polynomials and all interactions up to 4th degree (AKA 3rd order)
- F 6° polynomials and all interactions up to 11th order

[JMP col validation... fit all models with validation ... save prediction formula cols... analyze model... model comparison... complexity tradeoff illustration]
Conclusions? Model C looks the best. Note: another popular assessment metric besides oosRMSE is oosAAE which is just average absolute value difference. Strange ... given that linear models optimize for squared error. Show Demo with MLR w/ 10-fold CV *What **precisely** did I do that wasn't legal?*

A Possible Spin on Validation

Recall the proposal from last class:

- ➊ Split dataframe into training and test.
- ➋ Build model A on training.
- ➌ Predict using the test set.
- ➍ Calculate estimate of future generalization error of model 1.
- ➎ Build a different model B on training.
- ➏ Predict using the test set.
- ➐ Calculate estimate of future generalization error of model 2.
- ➑ ... steps 5-7 for model 3
- ➒ ... steps 5-7 for model 4
- ➓ ...
- ➑ ... steps 5-7 for model M
- ➒ Pick whichever model has better generalization error.

This is a form of **Model Selection**. What was wrong with is?

Looking into the Future is Not Legal

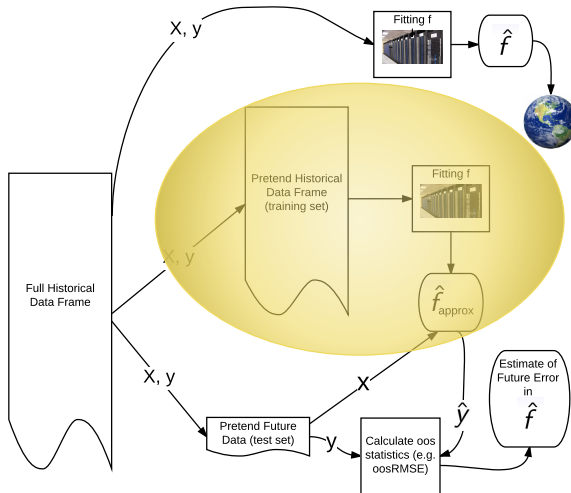
The oos validation is only valid if...



you treat the test set as a lockbox. Once you open it up, that's it!
And we opened it up M times!

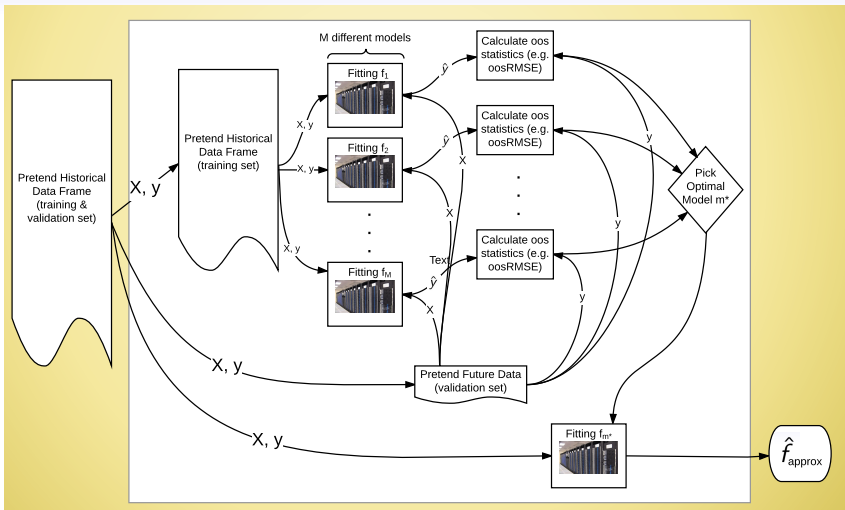
This is indeed a **Model Selection** procedure but ... our estimate of future generalization error is invalid. How can we do both?

In One Fold Let's Focus on the Training Set

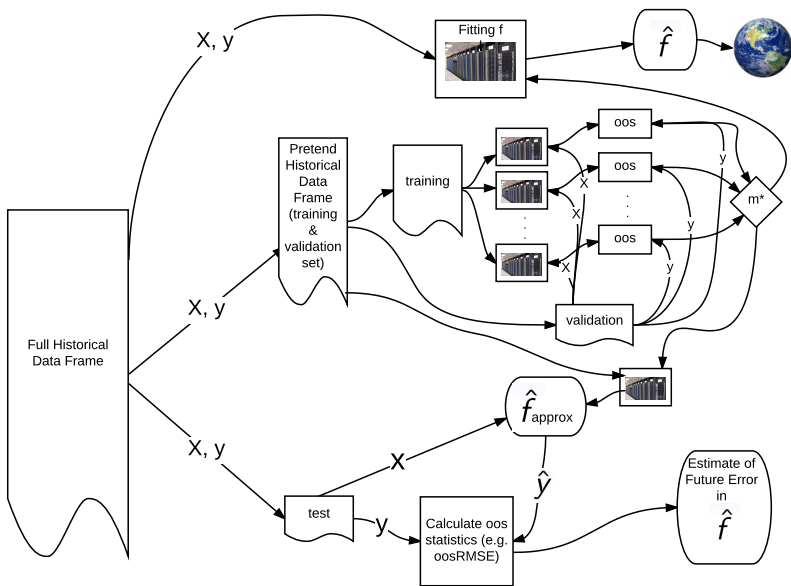


This procedure was completely valid as long as we did not touch the test set, right? As long as we operate only within the training set... we're OK!

Training \Rightarrow Training & Validation



3-way Splitting: The Full Picture



Why Training-Validation-Test Splitting?

- Training set: provides fits for many models where overfitting is “okay”
- Validation set: provides out-of-sample validation for each of the models. If the models are overfit, they will get wrecked at this stage.
- Test set: this lockbox provides a layer of security against overfitting within the training-validation union set.

Just like in the previous 2-way training-test split, you can overfit the training, get killed on the test set and be stuck. How could you similarly overfit here? Be careful of optimizing to the validation set. Models $1, 2, \dots, M$ should still be reasonable thought-through models.

Sizes of Training-Validation-Test Splits

Previously, we had the test set being 10-30%. This recommendation remains. Thus, the training-validation sets together should make up 70-90% of which a portion is the validation set (which is like a test set) and should be 10-30% of the total. Thus we arrive at proportions like 50-25-25 or 50-30-20 or 70-20-10. There is no exact guidance here.

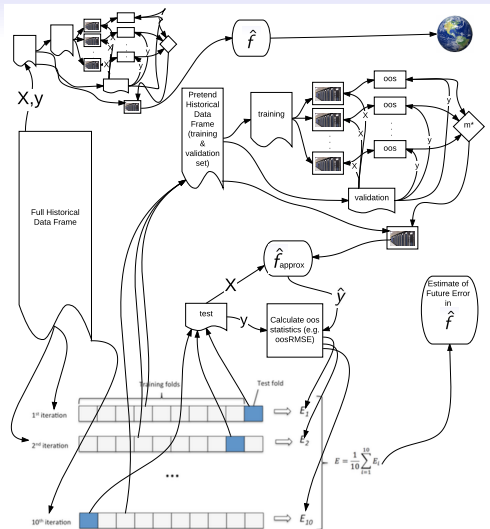
The same tradeoffs apply to the test set size but now we have new tradeoffs for the training set size vs. the validation set size:

- The larger the training set, the better the fit of the model (less bias) but the more variance in its assessment versus its peers
- The smaller the training set, the worse the fit of the model (more bias) but the less variance in its assessment versus its peers

Back to Wine...

We can do a single 3-partition split in JMP using the validation column... then data filter... then fit the 6 models again... then use model comparison to select best model ... then undo the filter ... then use model comparison again to find our test set error (the guess of the generalization error)... then build the full model for public consumption.

Can you CV this 3-Split Procedure?



Yes. It is called “nested resampling” and $K = 10$ fold CV is illustrated here. But... what are you now evaluating?

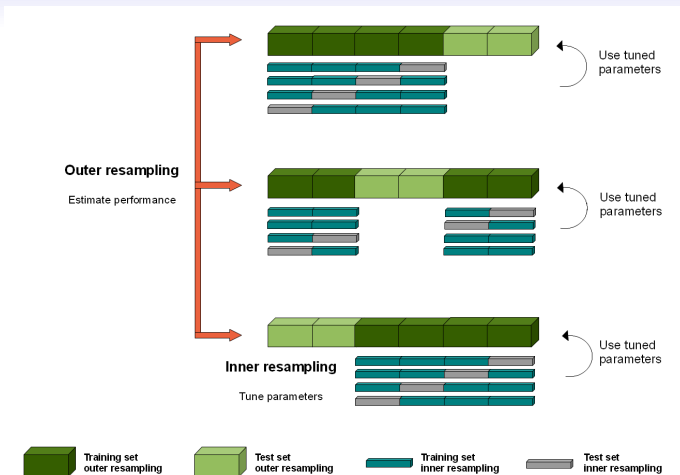
Without CV, it was just model m^* . But here m^* varies with the 10 folds! You are testing the entire procedure i.e. given models $1, 2, \dots, M$, pick the best one and ship it. How well you do in the future is estimated by the CV test set.

Not the First Means for Model Selection

CV with nested resampling is not generally done in the way it was described here as a means to evaluate a model selection procedure. Beyond scope of course: it is usually done to compare tuning settings in a non-parametric machine learning algorithm. We will see what this means next class.

Neither R nor JMP (to my knowledge) can do this out-of-the-box. In R, even with MLR, you have to program it (MLR uses it for tuning an algorithm). Maybe I will write to them?


Nested Resampling for Tuning



(from MLR's tutorial website).

3-way splits for model selection & evaluation

Forget CV for a moment since it complicates things ... the “inner split” consisted of the training-validation. We used this to “select” a model based on lowest oos error (RMSE, or highest R^2).



What's the cost? We had to use only 70-90% of the data (save the test set) to build the model. But what's the bigger cost?? You had to provide models 1, 2, ..., M!! This is no trivial task. So far I've been making them up! I've been using elaborate linear models which pivoted from parametric to non-parametric with the addition of curves and interactions. Forget finding the best model... how do we know we even employed good candidates?? **We don't.**